

# Mnist Case Study

Name: Youseef Osama Ahmed  
Name: Mohamed Al-Amin aba-Alziz

ID:20190629  
ID:20190720

In this case study we loaded the popular mnist dataset and tried different techniques, archs, activations ...etc. in order to find the best model to capture the complexity of the problem.

## Base Model:-

Model1:

final train: 0.9217 / Val: 0.9180

epoch 1: 0.4910 0.7171

epoch 2: 0.7866 0.8358

epoch 3: 0.8531 0.8727

epoch 4: 0.8781 0.8893

epoch 5: 0.8904 0.8981

#params 149834

Avg epoch time: 9s ~ 10s

Model Architecture:

Conv2D: 64 each(5, 5), strides=(2,2), activation=relu

MaxPooling2D: pool\_size(2, 2), strides(2, 2)

Dense: 64, activation='relu'

Dense: 10, activation='softmax'

Optimizers: SGD, lr 0.0001, momentum 0.9, epochs 10, batch size 32

## Trying different number of epochs

Model 2:

final train: 0.9325 0.9370

epoch 1: 0.4889 0.6504

epoch 2: 0.7332 0.8063

epoch 3: 0.8338 0.8562

epoch 4: 0.8679 0.8814

epoch 5: 0.8851 0.8920

#params 149834

Avg epoch time 9s ~ 10s

model Arch: Same as Model 1

Optimizers: epochs 15, ... SAME

# Increasing the epochs led to increase in the accuracy such that the model took extra time and training steps.

### Model 3

final train: 0.9412 0.9444

epoch 1: 0.4906 0.7059

epoch 2: 0.7799 0.8276

epoch 3: 0.8487 0.8644

epoch 4: 0.8740 0.8849

epoch 5: 0.8871 0.8941

#params 149834

Avg epoch time 9s ~ 10s

model Arch Same as Model 1

Optimizers epochs 20, ...SAME

# also increasing it more lead to more improvement but we prefer to step here in order to stop the model from overfitting

### Trying different learning\_rate

#### model 4:

final train: 0.9959 0.9848

epoch 1: 0.9480 0.9790

epoch 2: 0.9821 0.9837

epoch 3: 0.9864 0.9844

epoch 4: 0.9898 0.9856

epoch 5: 0.9921 0.9863

#params 149834

Avg epoch time 9s ~ 10s

model Arch Same as Model1

Optimizers lr 0.05, ...SAME

#Increasing the LR helped the model to converge faster

#### model 5:

final train: 0.9893 0.9861

epoch 1: 0.8103 0.9116

epoch 2: 0.9230 0.9420

epoch 3: 0.9458 0.9568

epoch 4: 0.9580 0.9644

epoch 5: 0.9657 0.9699

#params 149834

Avg epoch time 9s ~ 10s

model Arch Same as Model1

Optimizers lr 0.001, ...SAME

# the model improved slower than model 4

#### model 6:

final train: 0.9999 0.9906

epoch 1: 0.9247 0.9671

epoch 2: 0.9778 0.9822

epoch 3: 0.9842 0.9839  
epoch 4: 0.9874 0.9863  
epoch 5: 0.9910 0.9881  
#params 149834  
Avg epoch time 9s ~ 10s  
model Arch Same as Model1  
Optimizers lr 0.01, ...SAME  
#The model improvement is in between.

## Trying different architectures

model 7:

final train: 0.1000 0.9912  
epoch 1: 0.9461 0.9791  
epoch 2: 0.9830 0.9864  
epoch 3: 0.9885 0.9872  
epoch 4: 0.9908 0.9877  
epoch 5: 0.9927 0.9884  
#params 63242  
Avg epoch time 27s ~ 29s  
model Arch  
Conv2D: 32 each(3, 3), relu  
MaxPooling2D: pool\_size(2, 2), strides(2, 2)  
Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size(2, 2), strides(2, 2)  
Dense: 64, relu  
Dense: 32, relu  
Dense: 10, softmax

Optimizers SAME

#adding extra Conv layer improved the model since the model now  
captures extra patterns

model 8:

final train: 0.9837 0.9710  
epoch 1: 0.8571 0.9232  
epoch 2: 0.9439 0.9447  
epoch 3: 0.9549 0.9601  
epoch 4: 0.9627 0.9572  
epoch 5: 0.9668 0.9563  
#params 9322  
Avg epoch time 5s ~ 6s  
model Arch:  
Conv2D: 16, each(3, 3), strides(2, 2), relu  
MaxPooling2D: pool\_size(2, 2), strides=(2, 2)  
Conv2D: 32, each(3, 3), strides(2, 2), relu

MaxPooling2D: pool\_size(2, 2), strides(2, 2)

Dense: 64, relu

Dense: 32, relu

Dense: 10, softmax

Optimizers SAME

#decreasing the number of filters in the first Conv layer didn't hurt the model that much since it just learning basic features

model 9:

final train: 1.0000 0.9914

epoch 1: 0.9472 0.9806

epoch 2: 0.9829 0.9866

epoch 3: 0.9886 0.9822

epoch 4: 0.9910 0.9882

epoch 5: 0.9928 0.9883

#params 42698

Avg epoch time 30s ~ 31s

model Arch:

Conv2D: 32, each(3, 3), relu

MaxPooling2D: pool\_size(2, 2), strides(2, 2)

Conv2D: 32, each(5, 5), relu

MaxPooling2D: pool\_size(2, 2), strides(2, 2)

Dense: 32, relu

Dense: 10, softmax

Optimizers SAME

#increasing the strides improved the model since it now look at bigger area helping him to capture extra features

model 10:

final train: 1.0000 0.9881

epoch 1: 0.9282 0.9733

epoch 2: 0.9785 0.9754

epoch 3: 0.9852 0.9846

epoch 4: 0.9892 0.9847

epoch 5: 0.9928 0.9844

#params 693866

Avg epoch time 38s ~ 41s

model Arch:

Conv2D: 128, each(3, 3), relu

MaxPooling2D: pool\_size(2, 2), strides(2, 2)

Dense: 32, relu

Dense: 10, softmax

Optimizers: Same

#The model not as model 9 but 9 is better since 9 in just adding extra feature maps in the first Conv layers which doesn't help him but learning extra basic features not the complex ones

## Trying different batch sizes

model 11:

final train: 0.9992 0.9902

epoch 1: 0.9581 0.9838

epoch 2: 0.9843 0.9880

epoch 3: 0.9893 0.9864

epoch 4: 0.9916 0.9888

epoch 5: 0.9935 0.9899

#params 42698

Avg epoch time 28s ~ 29s

model Arch SAME as Model9

Optimizers Batch 16, ...Same

#Same no improvements

model 12:

final train: 0.9993 0.9897

epoch 1: 0.9231 0.9736

epoch 2: 0.9766 0.9814

epoch 3: 0.9835 0.9842

epoch 4: 0.9870 0.9861

epoch 5: 0.9887 0.9859

#params 42698

Avg epoch time 20s ~ 21s

model Arch Same as Model 9

Optimizers Batch 64, ...SAME

#Same no improvements

## Trying different activation functions

model 13:

final train: 0.9990 0.9906

epoch 1: 0.9434 0.9766

epoch 2: 0.9820 0.9850

epoch 3: 0.9880 0.9846

epoch 4: 0.9911 0.9860

epoch 5: 0.9934 0.9882

#params 63,242

Avg epoch time 19s ~ 20s

model Arch:

Conv2D: 32, each(3, 3), tanh

MaxPooling2D: pool\_size(2, 2), strides(2, 2)

Conv2D: 32, each(3, 3), tanh

MaxPooling2D: pool\_size(2, 2), strides(2, 2)

Dense: 64, tanh  
Dense: 32, tanh  
Dense: 10, softmax  
Optimizers: Same opt3  
#Same no improvements

model 14:

final train:

epoch 1: 0.9560 0.9771

epoch 2: 0.9835 0.9836

epoch 3: 0.9883 0.9876

epoch 4: 0.9911 0.9847

epoch 5: 0.9932 0.9879

#params 63,242

Avg epoch time 20s ~ 21s

model Arch: Same as Model 13 but wit SELU

Optimizers: Same opt3

# selu hurt the training it is always better to use the tan or relu in Conv

model 15:

final train: 0.9977 0.9887

epoch 1: 0.9432 0.9802

epoch 2: 0.9843 0.9853

epoch 3: 0.9889 0.9880

epoch 4: 0.9915 0.9878

epoch 5: 0.9932 0.9897

#params 63,242

Avg epoch time 20s ~ 22s

model Arch Same as model 13 but with LeakyRelu

Optimizers Same opt3

#Same no improvements

## Trying different optimizers

model 16:

final train: 0.9822 0.9749

epoch 1: 0.9517 0.9714

epoch 2: 0.9734 0.9708

epoch 3: 0.9768 0.9772

epoch 4: 0.9788 0.9802

epoch 5: 0.9810 0.9790

#params 42698

Avg epoch time 23s ~ 24s

model Arch Same as Model 9

Optimizers opt4

model 17:  
final train: 0.9812 0.9702  
epoch 1: 0.9460 0.9650  
epoch 2: 0.9742 0.9722  
epoch 3: 0.9768 0.9730  
epoch 4: 0.9781 0.9790  
epoch 5: 0.9776 0.9804  
#params 42698  
Avg epoch time 23s ~ 24s  
model Arch Same as model 9  
Optimizers opt5

## Trying different dropout rates

model 18:  
final train: 0.9973 0.9908  
epoch 1: 0.9424 0.9779  
epoch 2: 0.9828 0.9860  
epoch 3: 0.9874 0.9876  
epoch 4: 0.9905 0.9868  
epoch 5: 0.9925 0.9904  
#params 63242  
Avg epoch time 19s ~ 20s  
model Arch:  
Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size=(2, 2), strides=(2, 2)  
Dropout: 0.4  
Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size(2, 2), strides(2, 2)  
Dense: 64, relu  
Dense: 32, relu  
Dropout: 0.4  
Dense: 10, softmax  
Optimizers: SGD, lr 0.01, momentum 0.9  
#Same no improvements

model 19:  
final train: 0.9600 0.9851  
epoch 1: 0.7766 0.9693  
epoch 2: 0.9139 0.9766  
epoch 3: 0.9340 0.9786  
epoch 4: 0.9419 0.9846  
epoch 5: 0.9482 0.9843  
#params 63242

Avg epoch time 21s ~ 20s

model Arch:

Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size=(2, 2), strides=(2, 2)  
Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size=(2, 2), strides=(2, 2)  
Dropout: 0.4  
Dense: 64, relu  
Dropout: 0.4  
Dense: 32, relu  
Dropout: 0.4  
Dense: 10, softmax

Optimizers same as Model 18

#Same no improvements

model 20:

final train: 0.1120 0.1101  
epoch 1: 0.1112 0.1101  
epoch 2: 0.1110 0.1101  
epoch 3: 0.1117 0.1101  
epoch 4: 0.1111 0.1101  
epoch 5: 0.1110 0.1101

#params 63242

Avg epoch time 21s ~ 20s

model Arch:

Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size=(2, 2), strides=(2, 2)  
Dropout: 0.75  
Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size=(2, 2), strides=(2, 2)  
Dense: 64, relu  
Dense: 32, relu  
Dropout: 0.75  
Dense: 10, softmax

Optimizers same as Model 18

model 21:

final train: 0.4919 0.7359  
epoch 1: 0.1491 0.2524  
epoch 2: 0.2285 0.3947  
epoch 3: 0.3184 0.5078  
epoch 4: 0.3900 0.6452  
epoch 5: 0.4278 0.6483

#params 63242

Avg epoch time 19s ~ 20s

model Arch:



Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size=(2, 2), strides=(2, 2)  
Conv2D: 32, each(3, 3), relu  
MaxPooling2D: pool\_size=(2, 2), strides=(2, 2)  
Dropout: 0.75  
Dense: 64, relu  
Dropout: 0.75  
Dense: 32, relu  
Dropout: 0.75  
Dense: 10, softmax

Optimizers same as Model 18

#Increasing the dropout prevent the model from learning since most of the neurons will be now off

-----  
Summary of our best model, model 9

number of epochs is 20

size of batch is 32

stochastic SGD with a learning rate of 0.1 and a momentum of 0.9


we used ReLU as our activation function in all the layers except for the output layer where we used a softmax activation function

our architecture is as follow:

our first conv layer was 32 filters of size 3x3 followed by a max pool layer

the second conv layer is 32 filters of size 5x5 also followed by a max pool layer

We chose 1 FC layer consisting of 32 neurons followed by an output layer.



For extra info: <https://github.com/AI-ameen007/CNN/blob/main/CNN.ipynb>