

Mnist Case Study

Name: Youseef Osama Ahmed
Name: Mohamed Al-Amin aba-Alziz

ID:20190629
ID:20190720

In this case study we loaded the popular mnist dataset and tried different techniques, archs, activations ...etc. in order to find the best model to capture the complexity of the problem.

Optimizers: -

```
opt = SGD(learning_rate=0.0001, momentum=0.9)
opt1 = SGD(learning_rate=0.05, momentum=0.9)
opt2 = SGD(learning_rate=0.001, momentum=0.9)
opt3 = SGD(learning_rate=0.01, momentum=0.9)
opt4 = Adam(learning_rate=0.01, beta_1=0.9, beta_2=0.999)
opt5 = RMSprop(learning_rate=0.01, rho=0.9, momentum=0.1)
```

Base Model:-

Model1:

final train: 0.9217 / Val: 0.9180

epoch 1: 0.4910 0.7171

epoch 2: 0.7866 0.8358

epoch 3: 0.8531 0.8727

epoch 4: 0.8781 0.8893

epoch 5: 0.8904 0.8981

#params 149834

Avg epoch time: 9s ~ 10s

Model Architecture:

Conv2D: 64 each(5, 5), strides=(2,2), activation=relu

MaxPooling2D: pool_size(2, 2), strides(2, 2)

Dense: 64, activation='relu'

Dense: 10, activation='softmax'

Optimizers: SGD, lr 0.0001, momentum 0.9, epochs 10, batch size 32

Trying different number of epochs

Model 2:

final train: 0.9325 0.9370

epoch 1: 0.4889 0.6504

epoch 2: 0.7332 0.8063

epoch 3: 0.8338 0.8562

epoch 4: 0.8679 0.8814

epoch 5: 0.8851 0.8920

#params 149834

Avg epoch time 9s ~ 10s
model Arch: Same as Model 1
Optimizers: epochs 15, ... SAME
Increasing the epochs led to increase in the accuracy such that the model took extra time and training steps.

Model 3

final train: 0.9412 0.9444
epoch 1: 0.4906 0.7059
epoch 2: 0.7799 0.8276
epoch 3: 0.8487 0.8644
epoch 4: 0.8740 0.8849
epoch 5: 0.8871 0.8941
#params 149834
Avg epoch time 9s ~ 10s
model Arch Same as Model 1
Optimizers epochs 20, ...SAME
also increasing it more lead to more improvement but we prefer to step here in order to stop the model from overfitting

Trying different learning_rate

model 4:

final train: 0.9959 0.9848
epoch 1: 0.9480 0.9790
epoch 2: 0.9821 0.9837
epoch 3: 0.9864 0.9844
epoch 4: 0.9898 0.9856
epoch 5: 0.9921 0.9863
#params 149834
Avg epoch time 9s ~ 10s
model Arch Same as Mode1
Optimizers lr 0.05, ...SAME
#Increasing the LR helped the model to converge faster

model 5:

final train: 0.9893 0.9861
epoch 1: 0.8103 0.9116
epoch 2: 0.9230 0.9420
epoch 3: 0.9458 0.9568
epoch 4: 0.9580 0.9644
epoch 5: 0.9657 0.9699
#params 149834
Avg epoch time 9s ~ 10s

model Arch Same as Model1
Optimizers lr 0.001, ...SAME
the model improved slower than model 4
model 6:
final train: 0.9999 0.9906
epoch 1: 0.9247 0.9671
epoch 2: 0.9778 0.9822
epoch 3: 0.9842 0.9839
epoch 4: 0.9874 0.9863
epoch 5: 0.9910 0.9881
#params 149834
Avg epoch time 9s ~ 10s
model Arch Same as Model1
Optimizers lr 0.01, ...SAME
#The model improvement is in between.

Trying different architectures

model 7:
final train: 0.1000 0.9912
epoch 1: 0.9461 0.9791
epoch 2: 0.9830 0.9864
epoch 3: 0.9885 0.9872
epoch 4: 0.9908 0.9877
epoch 5: 0.9927 0.9884
#params 63242
Avg epoch time 27s ~ 29s
model Arch
Conv2D: 32 each(3, 3), relu
MaxPooling2D: pool_size(2, 2), strides(2, 2)
Conv2D: 32, each(3, 3), relu
MaxPooling2D: pool_size(2, 2), strides(2, 2)
Dense: 64, relu
Dense: 32, relu
Dense: 10, softmax
Optimizers SAME
#adding extra Conv layer improved the model since the model now
captuer extra patterns
model 8:
final train: 0.9837 0.9710
epoch 1: 0.8571 0.9232
epoch 2: 0.9439 0.9447
epoch 3: 0.9549 0.9601
epoch 4: 0.9627 0.9572

epoch 5: 0.9668 0.9563

#params 9322

Avg epoch time 5s ~ 6s

model Arch:

Conv2D: 16, each(3, 3), strides(2, 2), relu

MaxPooling2D: pool_size(2, 2), strides=(2, 2)

Conv2D: 32, each(3, 3), strides(2, 2), relu

MaxPooling2D: pool_size(2, 2), strides(2, 2)

Dense: 64, relu

Dense: 32, relu

Dense: 10, softmax

Optimizers SAME

#decreasing the number of filters in the first Conv layer didn't hurt the model that much since it just learning basic features

model 9:

final train: 1.0000 0.9914

epoch 1: 0.9472 0.9806

epoch 2: 0.9829 0.9866

epoch 3: 0.9886 0.9822

epoch 4: 0.9910 0.9882

epoch 5: 0.9928 0.9883

#params 42698

Avg epoch time 30s ~ 31s

model Arch:

Conv2D: 32, each(3, 3), relu

MaxPooling2D: pool_size(2, 2), strides(2, 2)

Conv2D: 32, each(5, 5), relu

MaxPooling2D: pool_size(2, 2), strides(2, 2)

Dense: 32, relu

Dense: 10, softmax

Optimizers SAME

#increasing the strides improved the model since it now look at bigger area helping him to capture extra features

model 10:

final train: 1.0000 0.9881

epoch 1: 0.9282 0.9733

epoch 2: 0.9785 0.9754

epoch 3: 0.9852 0.9846

epoch 4: 0.9892 0.9847

epoch 5: 0.9928 0.9844

#params 693866

Avg epoch time 38s ~ 41s

model Arch:

Conv2D: 128, each(3, 3), relu

MaxPooling2D: pool_size(2, 2), strides(2, 2)

Dense: 32, relu

Dense: 10, softmax

Optimizers: Same

#The model not as model 9 but 9 is better since 9 in just adding extra feature maps in the first Conv layers which doesn't help him but learning extra basic features not the complex ones

Trying different batch sizes

model 11:

final train: 0.9992 0.9902

epoch 1: 0.9581 0.9838

epoch 2: 0.9843 0.9880

epoch 3: 0.9893 0.9864

epoch 4: 0.9916 0.9888

epoch 5: 0.9935 0.9899

#params 42698

Avg epoch time 28s ~ 29s

model Arch SAME as Model9

Optimizers Batch 16, ...Same

#Same no improvements

model 12:

final train: 0.9993 0.9897

epoch 1: 0.9231 0.9736

epoch 2: 0.9766 0.9814

epoch 3: 0.9835 0.9842

epoch 4: 0.9870 0.9861

epoch 5: 0.9887 0.9859

#params 42698

Avg epoch time 20s ~ 21s

model Arch Same as Model 9

Optimizers Batch 64, ...SAME

#Same no improvements

Trying different activation functions

model 13:

final train: 0.9990 0.9906

epoch 1: 0.9434 0.9766

epoch 2: 0.9820 0.9850

epoch 3: 0.9880 0.9846

epoch 4: 0.9911 0.9860

epoch 5: 0.9934 0.9882

#params 63,242
Avg epoch time 19s ~ 20s
model Arch:
 Conv2D: 32, each(3, 3), tanh
 MaxPooling2D: pool_size(2, 2), strides(2, 2)
 Conv2D: 32, each(3, 3), tanh
 MaxPooling2D: pool_size(2, 2), strides(2, 2)
 Dense: 64, tanh
 Dense: 32, tanh
 Dense: 10, softmax
Optimizers: Same opt3
#Same no improvements

model 14:
 final train:
 epoch 1: 0.9560 0.9771
 epoch 2: 0.9835 0.9836
 epoch 3: 0.9883 0.9876
 epoch 4: 0.9911 0.9847
 epoch 5: 0.9932 0.9879
 #params 63,242
 Avg epoch time 20s ~ 21s
 model Arch: Same as Model 13 but wit SELU
 Optimizers: Same opt3
 # selu hurt the training it is always better to use the tan or relu in Conv

model 15:
 final train: 0.9977 0.9887
 epoch 1: 0.9432 0.9802
 epoch 2: 0.9843 0.9853
 epoch 3: 0.9889 0.9880
 epoch 4: 0.9915 0.9878
 epoch 5: 0.9932 0.9897
 #params 63,242
 Avg epoch time 20s ~ 22s
 model Arch Same as model 13 but with LeakyRelu
 Optimizers Same opt3
 #Same no improvements

Trying different optimizers

model 16:
 final train: 0.9822 0.9749
 epoch 1: 0.9517 0.9714
 epoch 2: 0.9734 0.9708

epoch 3: 0.9768 0.9772
epoch 4: 0.9788 0.9802
epoch 5: 0.9810 0.9790
#params 42698
Avg epoch time 23s ~ 24s
model Arch Same as Model 9
Optimizers opt4
model 17:
final train: 0.9812 0.9702
epoch 1: 0.9460 0.9650
epoch 2: 0.9742 0.9722
epoch 3: 0.9768 0.9730
epoch 4: 0.9781 0.9790
epoch 5: 0.9776 0.9804
#params 42698
Avg epoch time 23s ~ 24s
model Arch Same as model 9
Optimizers opt5
#No improvement opt5 is not better the extra Adam us helping

Trying different dropout rates

model 18:
final train: 0.9973 0.9908
epoch 1: 0.9424 0.9779
epoch 2: 0.9828 0.9860
epoch 3: 0.9874 0.9876
epoch 4: 0.9905 0.9868
epoch 5: 0.9925 0.9904
#params 63242
Avg epoch time 19s ~ 20s
model Arch:
Conv2D: 32, each(3, 3), relu
MaxPooling2D: pool_size=(2, 2), strides=(2, 2)
Dropout: 0.4
Conv2D: 32, each(3, 3), relu
MaxPooling2D: pool_size(2, 2), strides(2, 2)
Dense: 64, relu
Dense: 32, relu
Dropout: 0.4
Dense: 10, softmax
Optimizers: SGD, lr 0.01, momentum 0.9
#Same no improvements

model 19:

final train: 0.9600 0.9851

epoch 1: 0.7766 0.9693

epoch 2: 0.9139 0.9766

epoch 3: 0.9340 0.9786

epoch 4: 0.9419 0.9846

epoch 5: 0.9482 0.9843

#params 63242

Avg epoch time 21s ~ 20s

model Arch:

Conv2D: 32, each(3, 3), relu

MaxPooling2D: pool_size=(2, 2), strides=(2, 2)

Conv2D: 32, each(3, 3), relu

MaxPooling2D: pool_size=(2, 2), strides=(2, 2)

Dropout: 0.4

Dense: 64, relu

Dropout: 0.4

Dense: 32, relu

Dropout: 0.4

Dense: 10, softmax

Optimizers same as Model 18

#Same no improvements

model 20:

final train: 0.1120 0.1101

epoch 1: 0.1112 0.1101

epoch 2: 0.1110 0.1101

epoch 3: 0.1117 0.1101

epoch 4: 0.1111 0.1101

epoch 5: 0.1110 0.1101

#params 63242

Avg epoch time 21s ~ 20s

model Arch:

Conv2D: 32, each(3, 3), relu

MaxPooling2D: pool_size=(2, 2), strides=(2, 2)

Dropout: 0.75

Conv2D: 32, each(3, 3), relu

MaxPooling2D: pool_size=(2, 2), strides=(2, 2)

Dense: 64, relu

Dense: 32, relu

Dropout: 0.75

Dense: 10, softmax

Optimizers same as Model 18

model 21:

final train: 0.4919 0.7359

epoch 1: 0.1491 0.2524
epoch 2: 0.2285 0.3947
epoch 3: 0.3184 0.5078
epoch 4: 0.3900 0.6452
epoch 5: 0.4278 0.6483
#params 63242
Avg epoch time 19s ~ 20s
model Arch:

Conv2D: 32, each(3, 3), relu
MaxPooling2D: pool_size=(2, 2), strides=(2, 2)
Conv2D: 32, each(3, 3), relu
MaxPooling2D: pool_size=(2, 2), strides=(2, 2)
Dropout: 0.75
Dense: 64, relu
Dropout: 0.75
Dense: 32, relu
Dropout: 0.75
Dense: 10, softmax

Optimizers same as Model 18

#Increasing the dropout prevent the model from learning since most of the neurons will be now off

Summary of our best model, model 9:-

number of epochs is 20

size of batch is 32

stochastic SGD with a learning rate of 0.1 and a momentum of 0.9

we used ReLU as our activation function in all the layers except for the output layer


where we used a softmax activation function

our architecture is as follow:

our first conv layer was 32 filters of size 3x3 followed by a max pool layer

the second conv layer is 32 filters of size 5x5 also followed by a max pool layer

We chose 1 FC layer consisting of 32 neurons followed by an output layer.



For extra info: <https://github.com/Al-ameen007/CNN/blob/main/CNN.ipynb>